

# Lab 6: Reading and Writing Parquet Files in DBFS

**Author:** Dr. Sandeep Kumar Sharma

---

## Learning Objective

In this lab you will learn how to read from and write to Parquet files using Spark in Azure Databricks, and understand the benefits of Parquet over CSV/JSON for analytics workloads.

---

## Learning Outcome

By the end of this lab, you will be able to:

- Read CSV and write as Parquet
- Read Parquet files into DataFrames
- Understand partitioning and compression options
- Use Parquet for efficient querying and storage

---

## Lab Information

Source CSV (already available in DBFS):

```
/FileStore/tables/employee.csv
```

Parquet output directory:

```
/FileStore/tables/employee_parquet
```

---

## Step-by-Step Instructions

### Step 1 — Read the source CSV

```
input_csv = "/FileStore/tables/employee.csv"
df = spark.read.csv(input_csv, header=True, inferSchema=True)
display(df)
```

## Step 2 — Write DataFrame as Parquet (default options)

```
parquet_output = "/FileStore/tables/employee_parquet"
df.write.mode("overwrite").parquet(parquet_output)
# check files
display(dbutils.fs.ls(parquet_output))
```

## Step 3 — Read Parquet File

```
df_parquet = spark.read.parquet(parquet_output)
display(df_parquet)
```

## Step 4 — Use partitioning for large datasets (example)

If your data has a natural partition key (for example `department`), you can partition on write to speed up queries that filter by that column.

```
# Example: partition by department
df.write.mode("overwrite").partitionBy("department").parquet(parquet_output +
    "_partitioned")
# list partitioned output
display(dbutils.fs.ls(parquet_output + "_partitioned"))
```

## Step 5 — Compression options

Parquet supports compression codecs like `snappy`, `gzip`, `brotli`. Snappy is default and fast.

```
# write with snappy compression
(df.write
    .mode("overwrite")
    .option("compression", "snappy")
    .parquet(parquet_output + "_snappy"))
```

## Step 6 — Read with explicit schema (optional)

Providing a schema avoids `inferSchema` overhead on large files.

```
from pyspark.sql.types import StructType, StructField, IntegerType, StringType
schema = StructType([
    StructField("employee_id", IntegerType(), True),
```

```
StructField("name", StringType(), True),  
StructField("age", IntegerType(), True),  
StructField("department", StringType(), True)  
])  
  
# Read CSV with schema and write parquet  
  
df_with_schema = spark.read.csv(input_csv, header=True, schema=schema)  
df_with_schema.write.mode("overwrite").parquet(parquet_output + "_with_schema")
```

---

## Explanation & Best Practices

- Parquet is a columnar format which is efficient for analytic queries because it reads only the required columns.
  - Parquet stores schema with the data, so downstream readers automatically know column names and types.
  - Use partitioning when you have common filtering columns; it reduces the amount of data scanned.
  - Use compression (Snappy) for a good balance of speed and storage savings.
  - Prefer Parquet or Delta for production analytics instead of CSV for performance and metadata safety.
- 

## End of Lab 6

You have learned how to read and write Parquet files in DBFS, use partitioning, and apply compression. Next we can create a Delta lab, or a lab that demonstrates schema evolution and time travel with Delta.